CS Strategy: Razorpay Campus AI Build-a-thon

Student Brief and Problem Statements

1. At a glance

Item	Details
What you do	Pick one of three problems, build a working AI solution using the data we give you, and submit it with a 5-minute pitch video and short written answers.
Tools	Any. Claude, ChatGPT, Gemini, Copilot, n8n, Zapier, Google Sheets/Excel, Python, Lovable and similar. No coding is required, but you must be able to explain exactly what you built.
Data	Everything is synthetic (made-up). Use only the provided Data Pack.
Selection	Submissions are scored on a published rubric.

2. How the process works

1. **Pick one problem.** Read Section 4 and choose the problem you understand best. You can only submit for one.
2. **Build.** Use the Data Pack (Excel file) to build your solution. Focus on a working prototype, not slides alone.
3. **Submit.** Send three things: (a) your working solution (link, file, or screen recording), (b) a 5-minute pitch video, (c) your answers to the questions in Section 5.

3. Ground rules

- **Individual work.** You may discuss ideas, but the build and the pitch must be yours. You will be asked to explain every choice.
- **AI is expected, but you stay accountable.** Use AI tools freely. Tell us which ones, and what they did versus what you decided. If we spot a claim you cannot explain, it counts against you.
- **Synthetic data only.** Do not upload the Data Pack to tools you are not permitted to use, and do not add real merchant, customer or personal data.
- **The policies in the Data Pack are the source of truth.** They are fictional and may differ from real-world practice. Follow them as given.
- **No overclaiming.** Report results honestly, including what failed. A solution that is 70% accurate and honest beats one that claims 100% without proof.
- **Safety first.** Your solution must never invent facts, promise refunds, dates or waivers that the data does not support, or take irreversible actions without a human check.
- **Original work.** Do not copy another submission or published solution.

4. The three problems

Pick one. Each has a context, your task, the data you get, what we expect at minimum, and how we will measure it. Column details and sample rows are included here; the full data is in the Excel Data Pack.

Problem	In one line
Problem 1	Merchant support triage and resolution copilot: sort tickets, draft safe replies, escalate the risky ones.
Problem 2	Merchant churn early-warning and proactive outreach: find merchants likely to leave, say why, and decide what to do.
Problem 3	Merchant onboarding and KYC drop-off recovery: find why merchants stall and bring them back to activation.

Problem 1: Merchant support triage and resolution copilot

Context

Razorpay merchants are businesses that accept payments online. They write to support about settlements (money transfers to their bank), failed payments, refunds, chargebacks (customer disputes), KYC (identity verification), fees and API issues. Most tickets are repetitive but each needs the right data check and the right tone. A wrong or over-promising reply creates real risk.

Your task

Build an AI assistant that, for each ticket: (1) identifies the category, (2) decides what to do: auto-resolve, ask the merchant for more information, or escalate to a human team, (3) drafts a reply in the merchant's language style (English or Hinglish) using the merchant's status and the knowledge base, and (4) flags risk.

Data you get (tabs P1_Tickets, P1_Merchants, P1_KnowledgeBase)

- 150 merchant tickets from September 2026.
- 60 merchants with their current settlement status, hold reason, 7-day payment success rate, open chargebacks, pending refunds and KYC status.
- 17 knowledge-base articles with the policies and rules the support team must follow.

Tickets (sample)

ticket_id	created_at	channel	merchant_id	merchant_message
T-0001	01 Sep 2026 09:40	Email	M1043	Payout hold kyun kiya hai? Mera business ruk gaya hai, please release karo. M1043
T-0002	01 Sep 2026 13:26	Chat	M1015	Mere customers ka data leak hua hai tumhari wajah se. Main police me jaa raha hu. M1015
T-0003	01 Sep 2026 13:45	Email	M1034	I want a fee waiver for this month as my sales were low. Please approve. M1034
T-0004	01 Sep 2026 16:46	Chat	M1047	Many of my customers are saying UPI payments are failing at checkout since this morning. Please look into it. M1047

Merchants (sample)

merchant_id	segment	settlement_status	hold_reason	payment_success_rate_7d_pct	open_chargebacks	kyc_status
M1001	Mid	Settled	—	90.9	2	Verified
M1002	Small	Settled	—	92.9	1	Re-verification needed
M1003	Small	Delayed	Bank-side delay	97.4	1	Verified
M1004	Mid	Settled	—	91.1	2	Verified

Minimum expectations

- Run your solution on all 150 tickets and give us a table: ticket ID, category, action, escalation team (if any), risk level, drafted reply.
- Use the merchant's data. A reply about a delayed settlement must reflect that merchant's actual status.
- Show how you decide to escalate and what you do when information is missing.

Stretch ideas

- Detect frustrated or angry tone and adjust the reply. Spot a payment-failure spike across tickets. Suggest which knowledge-base article should be improved.

How we measure it

- Category accuracy and action accuracy against our hidden answers.
- Critical errors: invented dates, promised refunds or waivers, auto-replying to legal, fraud or security threats. These cost heavy points.
- Your own list of the replies you would not send, and why.

Problem 2: Merchant churn early-warning and proactive outreach

Context

Customer success managers (CSMs) often learn a merchant is leaving only after the volume has already moved to a competitor. The data to see it coming exists: falling payment volume (TPV, total payment value), dropping success rates, unresolved tickets, low usage and renewal dates. A CSM team can only contact a limited number of merchants each month, so the right merchants must be picked.

Your task

Build an AI agent that reads merchant signals and produces, for each merchant: a churn-risk score or label, the most likely reason in plain language, the recommended playbook, and a short drafted outreach message for the CSM to review. Then rank the merchants the team should contact first.

Data you get (tabs P2_Merchants, P2_Playbooks, P2_Assumptions)

- 100 merchants with six months of monthly TPV (Rs lakh, month 6 is most recent), success rates, support tickets, CSAT, products used, logins, contract renewal timing and an account manager note.
- Six CSM playbooks to choose from.
- Business assumptions: cost of a false alarm, value of a saved merchant, and a cap of 15 outreaches this month.
- **Not provided:** whether a merchant actually churned. We hold this label. Your job is to predict it.

Merchants (sample)

merchant_id	segment	tpv_m1_inr_lakh	tpv_m6_inr_lakh	success_rate_m1_m3_pct	success_rate_m4_m6_pct	unresolved_tickets	days_to_contract_renewal
C2001	Small	8.7	5.8	96.1	86.8	2	79
C2002	Small	2	1.9	94.6	94.1	0	278
C2003	Mid	59.9	82	94.9	95.3	0	356
C2004	Enterprise	275.8	228.5	94	93.4	0	39

Watch out: some merchants look risky but are not. Seasonal businesses dip after festive periods, tiny merchants have low volume by nature, and busy merchants raise many tickets because they are growing. Your solution should tell these apart from real risk.

Minimum expectations

- A ranked list of all 100 merchants with risk, reason and playbook; the top 15 flagged for outreach.
- Drafted outreach messages for the top 15, using real facts from each merchant's data.
- A short explanation of the logic or prompts that drive the ranking.

Stretch ideas

- Estimate revenue at risk using the assumptions. Group merchants by driver (performance, service, pricing, disengagement). Suggest what product or process fix would reduce the biggest driver.

How we measure it

- Precision and recall of your flagged merchants against hidden churn labels, within the 15-outreach cap.
- Whether your stated reasons match the true drivers, and whether decoys (seasonal, low-volume, busy-but-healthy) were wrongly flagged.
- Your estimate of the cost of false alarms versus value of merchants saved.

Problem 3: Merchant onboarding and KYC drop-off recovery

Context

Before a merchant can accept payments, they sign up, submit business details and KYC documents, verify a bank account, and pass a website or app review. Many stall: a document is rejected, a name does not match, or they simply stop. Each stalled merchant is revenue that never starts. Messages from reviewers are short, technical and not always clear to a small business owner.

Your task

Build an AI assistant that, for each stalled application: (1) diagnoses why it stalled, (2) explains the fix in simple language in the merchant's preferred language (English, Hindi, Hinglish, Tamil, Kannada or Marathi), (3) decides whether the case is recoverable or must be closed or routed to compliance, and (4) drafts the right nudge. Also tell the product team which fixable causes affect the most merchants.

Data you get (tabs P3_Applications, P3_RejectionCatalog)

- 100 onboarding applications with business type, language, step reached, days since last activity, attempts, support contacts, expected monthly TPV, and the reviewer's rejection message.
- A catalog of 13 reason codes with plain-language meaning, whether each is fixable, and the standard fix steps.
- **Note:** in about 15% of rejected cases the reason code is blank and only the reviewer's free text is available. Some applications have no rejection at all: the merchant just stopped.

Applications (sample)

application_id	business_type	preferred_language	step_reached	days_since_last_activity	rejection_code	rejection_message
A3001	Partnership	Marathi	KYC docs	38	R02	GST certificate could not be validated against GSTIN entered
A3004	Proprietorship	Hindi	KYC docs	3	R01	PAN ka naam aur application ka naam match nahi kar raha
A3005	LLP	English	KYC docs	19	—	—
A3006	Private Limited	Tamil	KYC docs	16	—	—

Minimum expectations

- For every non-activated application: diagnosed cause, fixable or not, recoverable or not, and a drafted nudge in the merchant's language and a suitable channel (WhatsApp only if opted in).
- Never suggest a workaround for restricted business categories, and never reveal compliance screening reasons to the merchant.
- A summary: how many applications are recoverable, expected TPV at stake, and the top three systemic causes.

Stretch ideas

- Prioritise nudges by TPV and days since activity. Propose a product change to prevent the top cause. Design a follow-up sequence with stopping rules.

How we measure it

- Diagnosis accuracy against hidden answers, including the blank-code cases.
- Correct handling of the non-fixable and compliance cases, and fixed-versus-recoverable counts.
- Clarity, tone and language quality of the nudges.

5. What to submit

A. Your working solution

Share a link or file we can open (spreadsheet, shared chat, workflow export, app link, or repository), plus the output table described in your problem. A screen recording of the solution running is acceptable.

B. Pitch video (max 5 minutes)

1. Problem and user (30 seconds): who has this problem and what does it cost them?
2. Your approach (1 minute): what does AI do, and what is left to rules or humans?
3. Live demo (1.5 minutes): run it on real rows from the data.
4. Results and failures (1.5 minutes): your numbers and where it went wrong.
5. Next steps (30 seconds): what you would do with three months and a team.

C. Written answers

Answer in a short document or form. Keep each answer to 100 to 150 words.

All problems

1. Which problem did you pick, and who is the main user (merchant, CSM, support agent)?
2. What does a good outcome look like in numbers today, and what does your solution change?
3. What does AI do in your solution, and what is deliberately left to rules or humans? Why?
4. Which tools and models did you use, and why those?
5. How did you test it? Show results on at least 20 cases, including the ones it got wrong.
6. Describe one real failure or edge case, what it could have cost, and how your design handles it.
7. What are the risks (wrong answer, bias, compliance, privacy), and what controls did you add?
8. If you had 3 months and a small team, what would you build next, and how would you know it worked?

Problem-specific (answer the set for your problem)

Problem	Extra questions
Problem 1: Triage copilot	How do you decide when to escalate? What share of tickets can be safely auto-resolved, and at what accuracy? What does your solution do when a merchant is angry or the message is ambiguous?
Problem 2: Churn early-warning	Which signals mattered most, and why? How did you avoid flagging seasonal and low-volume merchants? How should the playbook differ for a large merchant versus a small one?
Problem 3: Onboarding recovery	How do you tell a fixable rejection from a genuine compliance failure? How does your nudge change by language, channel and timing? Which single fix would you ask the product team to ship first?

6. How you will be evaluated

Criterion	Weight	What strong looks like
Problem understanding and user insight	20	Clear on who the user is, what hurts, and what success means in numbers.
Quality and working state of the solution	25	It runs on the full data, is easy to follow and does what it claims.
Measured results and honest failure analysis	20	Real metrics, a list of mistakes and what you learned from them.
Judgment on risk, guardrails and human-in-the-loop	15	Knows what AI should not decide and builds checks for it.

Criterion	Weight	What strong looks like
Pitch clarity and Q&A	20	Crisp, structured, confident answers; can defend every choice.

Deal-breakers: using real or confidential data, copying another submission, or presenting results you cannot reproduce. Any of these leads to disqualification.

7. Frequently asked questions

Do I need to know how to code?

No. Many strong submissions use no-code tools. You must be able to explain what you built and why.

Can I use the data in any AI tool?

Yes, it is synthetic. Do not add real merchant or personal data.

Can I work with a friend?

No. This is an individual challenge. You can discuss ideas, but the work is yours.

What if the data is not enough?

State your assumption clearly and move on. Good judgment on gaps is part of the evaluation.

What if my solution is not perfect?

That is fine. Be honest about what works and what does not. We reward honest results over polished claims.

Appendix: Glossary

Term	Meaning
Merchant	A business that accepts payments through Razorpay.
Settlement	Transfer of collected money from Razorpay to the merchant's bank account, usually T+2 business days after payment.
TPV	Total payment value: the rupee amount of payments a merchant processes. Here in Rs lakh (1 lakh = 100,000).
Success rate	Share of payment attempts that succeed.
MDR	Merchant discount rate: the percentage fee charged on payments.
Refund	Returning a customer's money to their original payment method.
Chargeback	A dispute raised by a customer with their bank; the merchant must submit evidence within a deadline.
KYC	Know Your Customer: identity and business verification using documents such as PAN, GST and bank proof.
UPI	Unified Payments Interface, India's instant bank-to-bank payment system.
CSM	Customer success manager: the person who manages a merchant relationship.
Churn	A merchant moving most of its payment volume away from Razorpay.
Escalate	Pass a case to a specialist team (for example Risk Ops, Legal, Payments Ops).